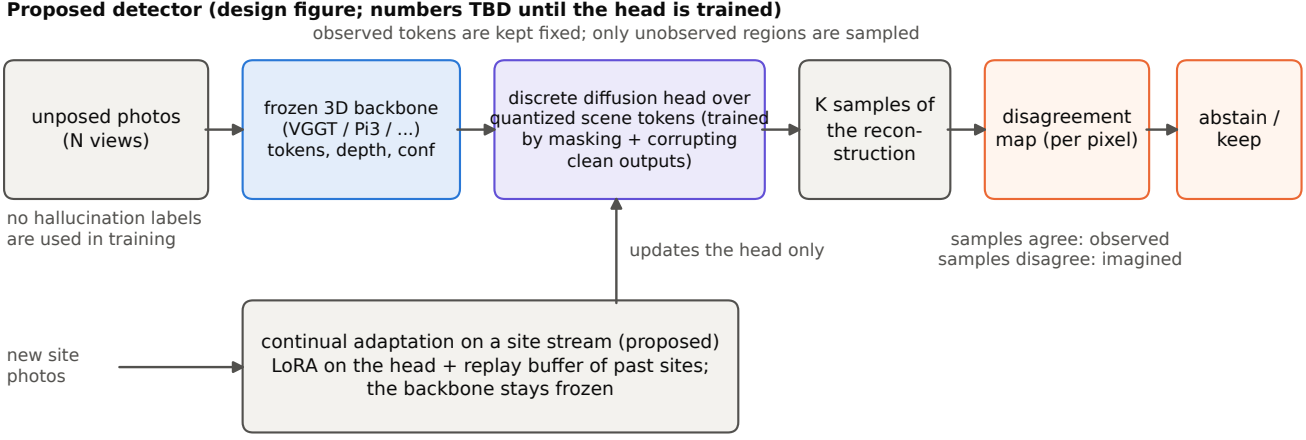
Proposed detector (design figure; numbers TBD until the head is trained)
observed tokens are kept fixed; only unobserved regions are sampled
unposed photos (N views)
frozen 3D backbone (VGGT / Pi3 / ...) tokens, depth, conf
discrete diffusion head over quantized scene tokens (trained by masking + corrupting clean outputs)
K samples of the recon-struction
disagreement map (per pixel)
abstain / keep
no hallucination labels are used in training
updates the head only
samples agree: observed
samples disagree: imagined
new site photos
continual adaptation on a site stream (proposed)
LoRA on the head + replay buffer of past sites;
the backbone stays frozen