Observe, Imagine, Abstain: Detecting Hallucinated Geometry in Feed-Forward 3D Reconstruction

Anonymous CVPR submission



Figure 1. **Native confidence reads the patch, not the missing support.** Two identical gray patches: red is hidden in every input view (imagined), green is hidden here but visible in another view (observed). Native confidence is dark on both; our label-free support head, which reads the backbone's cross-view tokens, is dark only where no view sees the surface. Right: AUROC ranges over models and datasets (runs P11-matched, P4-vis*, P7-vis-imc).

Abstract

*Feed-forward 3D models reconstruct a scene from unposed photographs in one pass. Where no photograph shows a surface they invent one, and report it with the confidence of an observed surface. We build HalluBench-3D, which hides a 3D region of the ground truth from every input view, and test five models on 18 DTU scans and 18 CO3Dv2 scenes: native confidence lets 24% to 81% of the invented geometry through a filter that keeps 90% of the observed geometry (blur and noise fills). A matched split shows why. Against an identical patch that another view does see, confidence is at chance (*AUROC *0.43 to 0.55): it reads the patch, not the missing support. The information is in the model's cross-view tokens. A small* support head *(1.7 to 3.3 M parameters), trained on the backbone's own multi-view agreement with no ground truth, flags single-view geometry at* AUROC *0.93 to 0.99 on four datasets it never saw (native confidence: 0.55 to 0.89), and separates the matched patches where a per-image control cannot. Against dense scans of real scenes it finds surfaces no input view saw better than native confidence in 10 of 14 pre-registered settings.*

1. Introduction

Feed-forward 3D reconstruction has moved from pairwise pointmaps [14, 31] to models that take dozens of unposed images and return dense geometry in one pass [13, 17, 29, 32]. They output geometry at every pixel, including pixels whose surface no input view shows: the objective rewards a plausible guess, so they learn to guess. A conservator who sees a crack in a reconstructed dome needs to know whether a photograph shows it or the model drew it from what domes look like. Today the confidence map is bright in both cases.

What we find.

1. **Confidence does not know what the model saw.** A quarter to four fifths of the invented geometry passes a confidence filter, and against an identical patch that another view sees, confidence is at chance.
2. **The cross-view tokens do.** A linear probe on them reaches AUROC 0.99; a head trained without labels reads the same signal on four backbones and four unseen datasets.
3. **A consistent invention passes every cross-view test,** ours included, and five models invent the same surface. It follows the pixels it was drawn from.

Contributions. (i) A definition of geometric hallucination and HalluBench-3D, with a *matched split* that removes the patch confound and makes a per-image control a required test for any detector of unobserved geometry (Sections 3 and 4). (ii) A label-free *support head* on the frozen backbone that abstains between observed and imagined (Figure 5, Section 5), tested on four backbones and four datasets against native confidence, reprojection, a per-image control and a ground-truth-trained head (Section 6). (iii) Controls

on identical rows, a real-occlusion check against a dense scan, and the limit of every cross-view test. Every number comes from a logged run with a run id, and every pass condition was registered, with a date, before the run it judges (supplement).

2. Related work

Feed-forward 3D. DUSt3R and MASt3R [14, 31] predict pointmaps for image pairs; VGGT [29], Pi3 and Pi3X [32], MapAnything [13], VGGT-Omega [30] and Depth Anything 3 [17] take many unposed views in one pass. Each exposes a per-pixel confidence trained as a loss weight on observed data; none is trained to tell observed from unobserved surfaces.

Hallucination in feed-forward 3D. Attention heads of DUSt3R, VGGT and DA3 keep matching correspondences after small patches are occluded [3], measured with attention and Sampson error rather than confidence. Unrelated scenes, repeated images and noise still produce dense geometry [23]. Monocular depth fails on illusions [21]; wide viewpoint spans produce unsupported geometry [15]; unobserved parts of a room escape inference-time uncertainty [5]; inpainting fills masks on purpose [22]. None has per-point labels. In vision-language models POPE [16] labels statements the image does not support; ours is the geometric analogue.

Uncertainty for feed-forward 3D. ReVar3R [18] takes per-point variance over perturbed runs and, by its authors' account, misses stable systematic bias, which a confidently filled region is. Trust3R [33] fits an evidential head, GeoCond [1] adds pose-level uncertainty, and Deng et al. [7] find the implied uncertainty of seven backbones at least 2.4× too small. All score uncertainty against error, which conflates confident errors on observed surfaces with confident predictions on unseen ones; CUR isolates the second. Closest to our detector, a linear probe on frozen VGGT features predicts camera-pose failure [28] and a small head predicts pairwise co-visibility [34]; our head is the dense, per-pixel analogue for unobserved *geometry*, trained with no label at all.

Generative 3D. Hidden surfaces are decodable from frozen VGGT features [2]; feed-forward splatting [9, 26] and world models [4] store imagined and observed surfaces alike; sample disagreement [11, 12, 20] and quantized generative heads [6, 19] are the natural successors to a discriminative head.

3. HalluBench-3D

Definition. A model returns depth D_i and confidence C_i for each input view i. A pixel is *supported* if its ground-truth surface point is visible in at least one input view and *unsupported* otherwise; a model hallucinates when it outputs confident geometry at an unsupported pixel. Native confidence is a loss weight learned on observed data, so a smooth invented surface is confident by construction.

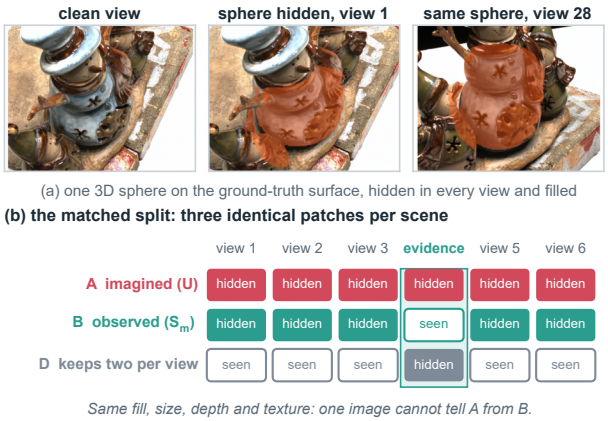


Figure 2. **The benchmark and its matched split.** (a) A sphere on the ground-truth surface is hidden in every view, so no camera sees it; its pixels are unsupported. (b) The matched split adds identical patches that one evidence view still sees, so only cross-view evidence separates imagined from observed.

Metrics. AUROC is $P(C_S > C_U)$, with 0.5 at chance. The Confident-Unsupported Rate CUR sets τ so that 90% of supported pixels pass and counts the unsupported pixels that pass too: the hallucination rate after filtering. CURERR adds a depth error above 10%. Means carry 95% cluster-bootstrap intervals over scenes; comparisons use paired sign-flip tests on per-scene means, Holm-corrected ($p < 0.0001$ is the test's floor).

Protocol. Ten views per scene. A sphere of radius 0.10 or 0.18 of the object diagonal is centred on the ground-truth surface and every pixel whose ground-truth point lies inside it is hidden in every view (Figure 2a) and filled with gray, a blur, noise or an inpainting [27]. Hidden pixels are U, pixels outside a 6-pixel margin are S. Five seeds per scene.

The matched split. Every U pixel sits under a patch and no S pixel does, so a confidence that merely reacts to a patch would look as if it knew about missing support. The matched split (Figure 2b) places three spheres matched in area, depth, texture and fill: A is hidden in every view (U); B is hidden except in an evidence set of one, two or half the views, so its pixels are filled but observed (S_m); D is hidden only in the evidence views, so every view carries two patches. AUROC_m and CUR_m score U against S_m.

Data. *DTU* [8]: the 18 scans of the MVSNet validation split, 10 views spread over the 49-camera rig. *CO3Dv2* [24]: 18 phone orbits, the two depth-carrying sequences of nine categories, with COLMAP depth. *ETH3D* [25]: nine building-scale DSLR scenes with laser-scanner ground

Table 1. **The leak, and what it is.** AUROC and CUR of native confidence for invented against clean observed pixels (blur and noise, 18 scenes × 5 seeds × 2 radii), and, shaded, against an identical patch that another view sees (matched split). Intervals in the supplement.

	DTU (lab rig)				CO3Dv2 (phone)			
	vs clean		vs observed patch		vs clean		vs observed patch	
model	AUROC↑	CUR↓	AUROC↑	CUR↓	AUROC↑	CUR↓	AUROC↑	CUR↓
VGGT-1B	0.90	0.24	0.55	0.80	0.68	0.65	0.49	0.85
Pi3	0.91	0.28	0.49	0.83	0.89	0.27	0.49	0.81
Pi3X	0.81	0.53	0.47	0.87	0.73	0.57	0.50	0.81
MapAnything	0.74	0.62	0.51	0.84	0.58	0.81	0.49	0.84
Depth Anything 3	0.71	0.70	0.43	0.85	0.62	0.80	0.48	0.84

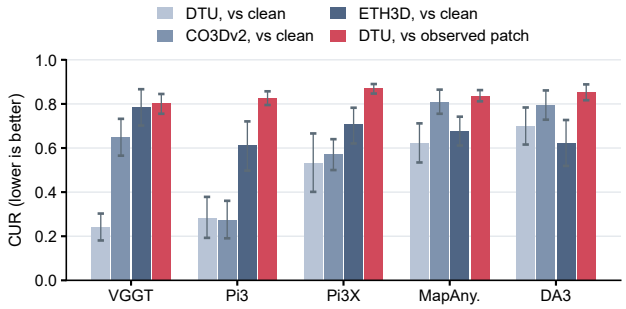


Figure 3. **Every model leaks, and more against the matched control.** CUR of native confidence (95% intervals over scenes) against clean pixels on three datasets, and against an identical observed patch (red). Lower is better.

truth. *Phototourism* [10]: nine landmarks of the Image Matching Challenge 2020, Internet photographs with people in front of the monument and unordered focal lengths. Every threshold is frozen on DTU. Three backbones list CO3Dv2 in their training data; none lists DTU or ETH3D.

Models. VGGT-1B [29], Pi3 and Pi3X [32], MapAnything [13] and Depth Anything 3 [17], each with its own depth and confidence, at the same views and resolution. A Gaussian-splatting model [9], scored with its opacity, gives the same answer (supplement).

4. Native confidence reads the patch

The leak. Table 1 and Figure 3: every model lets a large share of invented geometry through, from CUR 0.24 for VGGT to 0.70 for Depth Anything 3 on DTU. Three tiers emerge: VGGT and Pi3 drop their confidence inside the hidden region yet pass one in four unsupported pixels; Pi3X passes one in two; MapAnything and DA3 are close to blind. The gray fill, an obvious patch, gives the smallest numbers (down to 0.053). The invented depth is accurate (ABSREL_U 0.011 to 0.022): the failure is that imagined cannot be told from observed, not that the imagination is poor. MapAnything is the exception, with 29% of its unsupported pixels confident and wrong by more than 10%. After Holm correction VGGT and Pi3 are indistinguishable ($p = 0.41$) and both leak less than the other three ($p \leq 0.0015$).

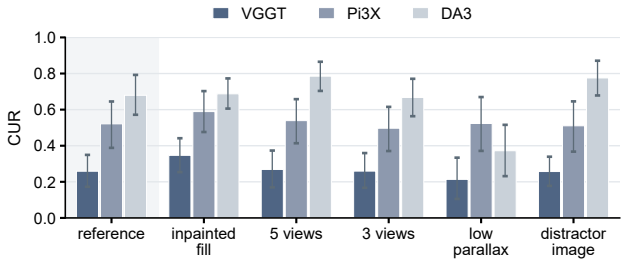


Figure 4. **Stress tests on DTU** (blur, radius 0.18, seeds 0 and 1). A seamless inpainted fill raises the leak; fewer views or a distractor image leave it unchanged. Runs P2-stress, P1-full.

Phone captures change the ranking. On CO3Dv2 Pi3 is the only model whose leak transfers from the rig (CUR 0.27 on both); VGGT, joint best on DTU, leaks 2.7× more, and MapAnything and DA3 pass four in five. The leak is largest on the dataset three of the models were trained on: seeing CO3Dv2 did not teach their confidence to notice a missing surface.

What the confidence reads. Shaded blocks: against an identical patch that another view does see, native confidence separates nothing. AUROC_m is 0.43 to 0.55 on DTU and 0.48 to 0.50 on CO3Dv2, and a threshold that keeps 90% of the observed patch keeps 80% to 87% of the imagined one. Clean pixels against the observed patch score 0.75 to 0.92: the original AUROC is the confidence dropping under a patch, whichever label it carries. The pre-registered rule (Gate 6) marks every model as patch-inflated on DTU, so the matched numbers are the ones to read.

Not the fill, not the views, not cheap fixes. An inpainted fill without seams raises the leak for VGGT (+0.088, $p = 0.031$) and Pi3X (+0.068, $p = 0.005$), so Table 1 is a lower bound for a natural-looking occluder (Figure 4). Dropping from 10 to 3 views leaves VGGT and Pi3X unchanged ($p > 0.4$): fewer views to cross-check do not make the models less sure, so native confidence follows single-view appearance. DA3 differs, its leak rising with a distractor and halving at low parallax, consistent with a confidence that tracks appearance consistency across views. Variance over four perturbed runs is far worse than native confidence (AUROC 0.67 and 0.55 against 0.92 and 0.83) and reprojection with the model's own cameras is at chance on the patches (0.49, 0.48): the invented surface is stable and cross-view consistent.

Takeaway. The shipped confidence detects the look of a patch. A detector of unobserved geometry must pass a test where appearance cannot explain the answer.

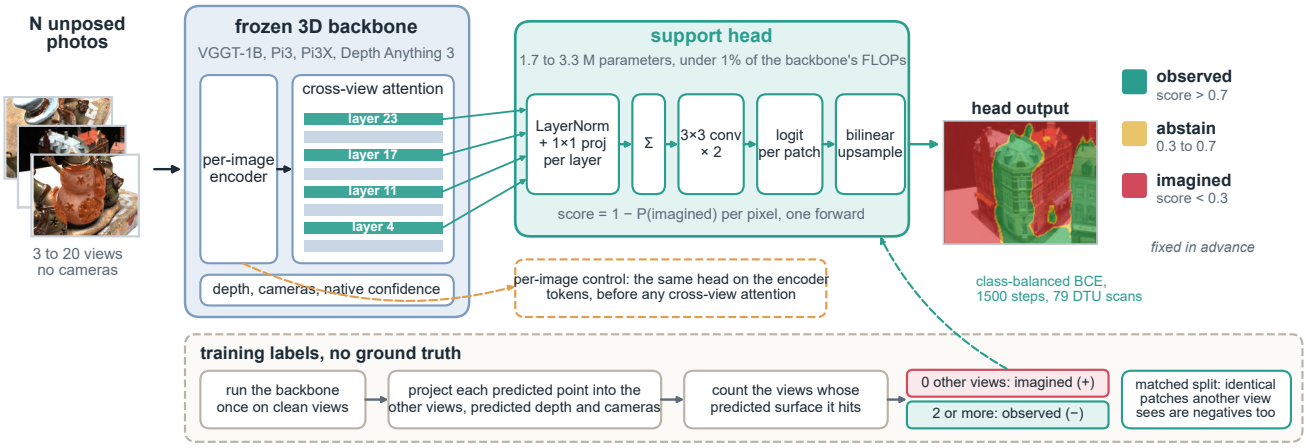


Figure 5. **The support head.** A frozen backbone's cross-view layers feed a small head that returns, per pixel, whether the output is corroborated by other views, and abstains in between. Its training labels come from the backbone's own prediction: each predicted point is reprojected into the other views and counted. No ground truth, no change to the backbone, one extra forward under 1% of its cost. The per-image control is the same head on tokens taken before any cross-view attention.

5. Method: a support head that abstains

Architecture. The backbone (VGGT, Pi3X, Pi3 or Depth Anything 3) stays frozen. The head taps the tokens its depth or confidence head reads (for VGGT and DA3 four cross-view layers, 2048-d at 14×14 patches), applies LayerNorm and a 1×1 projection per layer, sums, applies two 3×3 convolutions and returns one logit per patch (Figure 5). The score is $1 - P(\text{imagined})$; three bands fixed in advance tell the user *observed* (> 0.7), *abstain* or *imagined* (< 0.3).

Labels without ground truth. The backbone runs once on clean views; each predicted point is projected with the predicted depth and cameras into the other views and counted where it lands on their predicted surface. No other view: imagined. Two or more: observed. This is the reprojection test used as a teacher, so the head must generalise beyond it to be worth having. Ground-truth visibility gives the upper bound. *Matched labels* use the three-patch protocol, with the observed patches as negatives so that a patch cue is penalised. All heads: class-balanced binary cross-entropy, 1500 AdamW steps (learning rate 10^{-3}, cosine schedule), 3 to 10 random clean views per step on the 79 DTU training scans, disjoint from every benchmark scene; ten minutes on one H200, hyperparameters fixed before any benchmark scan was evaluated.

Controls. The decisive control is the *per-image head*: same head, same labels, on tokens before cross-view attention. Under matched labels it must stay near chance, or the pipeline leaks. Further: a linear probe per layer, image-border and depth-edge cues, extra seeds, and a provenance variant with cross-view attention that says *which* view corroborates a pixel (supplement).

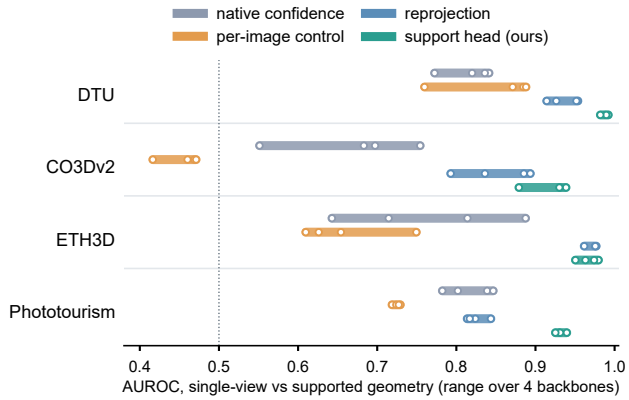


Figure 6. **Four datasets the head never saw.** AUROC for single-view against supported geometry on clean views (3, 5 and 10), range over four backbones (dots). The per-image control falls below chance off the lab rig. Full table in the supplement.

Table 2. **Single-view geometry on clean views,** mean over four backbones (VGGT-1B, Pi3X, Pi3, DA3; per-backbone numbers and intervals in the supplement). Our heads are tinted, the best value per column is bold. No head saw these scenes.

	DTU		CO3Dv2		ETH3D		Photot.	
detector	AUROC↑	CUR↓	AUROC↑	CUR↓	AUROC↑	CUR↓	AUROC↑	CUR↓
native confidence	0.82	0.49	0.67	0.80	0.76	0.50	0.82	0.46
image-border cue	0.70	0.73	0.45	0.90	0.73	0.63	0.73	0.65
per-image control	0.85	0.40	0.45	0.89	0.66	0.72	0.72	0.68
reprojection, own cameras	0.94	0.13	0.85	0.40	**0.97**	**0.07**	0.82	0.34
support head, GT labels	**0.99**	**0.02**	**0.92**	0.24	**0.97**	0.08	**0.93**	0.21
support head, label-free	**0.99**	0.03	**0.92**	**0.23**	**0.97**	0.09	**0.93**	**0.19**

6. Experiments

Four datasets the head never saw. Figure 6 and Table 2: trained on DTU training scans only, the label-free head separates single-view from supported geometry at AUROC 0.98

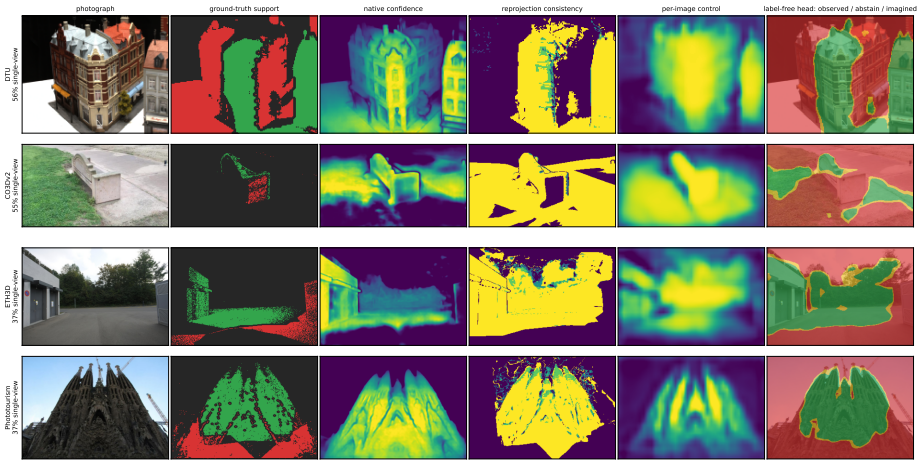


Figure 7. **What the label-free head says on unseen captures.** Per dataset: photograph; ground-truth support (red: no other view sees it, green: two or more do); native confidence; reprojection; per-image control; the head's bands (green observed, yellow abstain, red imagined). On DTU native confidence lets half of the single-view pixels through a 90% filter, the head 1.7%. Run P6-dump.

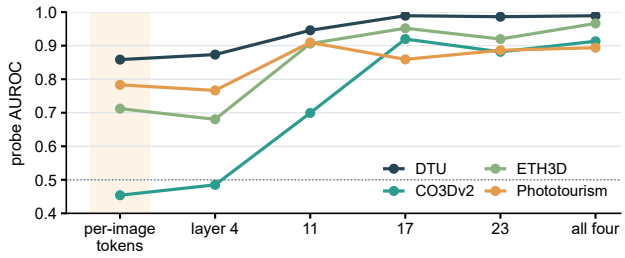


Figure 8. **Where the signal lives.** A linear probe on one VGGT layer at a time. Per-image tokens (shaded) carry only a frame-edge cue that fails off the rig; support becomes linearly decodable from layer 17 of 24.

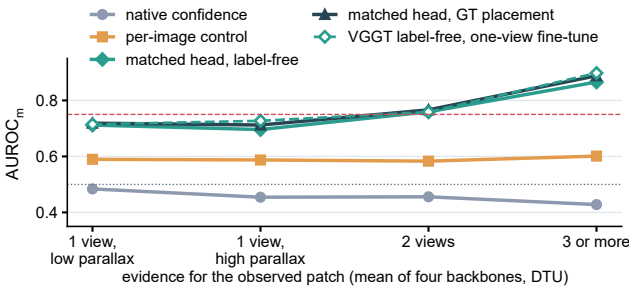


Figure 9. **The matched head reads evidence.** AUROC$_m$ of the imagined against the observed patch, by how many clean views see the observed one (mean of four backbones). Native confidence and the per-image control stay near chance; the head improves with evidence. Dashed: the pre-registered 0.75 and the one-view fine-tune.

to 0.99 on DTU, 0.88 to 0.94 on CO3Dv2, 0.95 to 0.98 on ETH3D and 0.92 to 0.94 on the landmark photographs, against 0.55 to 0.89 for native confidence. It is within 0.02 of a head trained on ground-truth visibility, beats its teacher, reprojection, on DTU, CO3Dv2 and the landmarks, and ties it on ETH3D, where adjacent DSLR views make reprojection as good as any head. Gate 3 asked the ground-truth head to beat native confidence and the per-image control by 0.02 AUROC on DTU and one unseen dataset; Gate 4 asked the same of the label-free head, plus being within 0.03 of the ground-truth head and ahead of its teacher on CO3Dv2. Both were written before training and both pass. On the landmark photographs, where people stand in front of the monument and the model's own cameras are least reliable, the head beats reprojection by 0.11 to 0.13. Figure 7 shows what this looks like.

Where the signal lives. The per-image control falls below chance on CO3Dv2 and ETH3D: without cross-view evidence the head learns the frame edge, which marks single-view pixels on the rig and nowhere else. A linear probe on one layer at a time (Figure 8) shows support becoming linearly decodable in the second half of the aggregator, and a probe on all four tapped layers is within 0.015 of any head we trained: the information is in the features, and the head is a reader. A head trained on the benchmark's fills instead learns the fill (AUROC 1.000, and so does its per-image control) and scores below native confidence on real single-view geometry: masking labels need the control.

The matched head. Trained on the three-patch protocol on all four backbones (Figure 9), the per-image control stays at 0.56 to 0.61 AUROC$_m$ (pre-registered bound 0.60), so the pipeline does not leak. The head separates the imagined patch from the observed one at AUROC$_m$ 0.69 to 0.77 against 0.43 to 0.55 for native confidence ($p \leq 0.0003$ on every backbone), the label-free head within 0.04 of the ground-truth one, and holds on CO3Dv2 without retraining

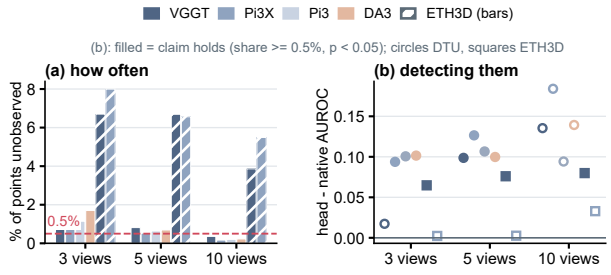


Figure 10. **Real occlusion against dense scans** (DTU, 49 views; ETH3D, every image of the scene). (a) Share of predicted points on surfaces no input view saw; a claim is allowed above 0.5% (dashed). (b) The head's AUROC gain over native confidence on those points; filled where the pre-registered claim holds.

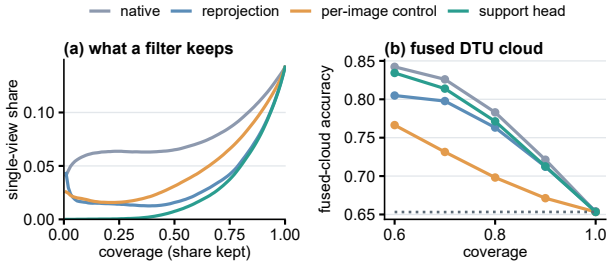


Figure 11. **What abstaining buys** (VGGT, DTU). (a) Single-view share of the pixels a filter keeps, by coverage. (b) Accuracy of the fused point cloud when the least trusted points are dropped. Native confidence tracks error, the head tracks support.

(0.62 to 0.69 against 0.54 to 0.58 for its control). Its accuracy rises with the evidence, from 0.68 to 0.748 with one clean view at low parallax to 0.80 to 0.92 with three or more: the signature of a detector that reads cross-view support, not appearance. A single low-parallax view is the hardest case, and longer training does not move it (Gate 9). The check earned its place: a first placement that ignored texture let the observed patch fall on the rig's backdrop and the control rose to 0.62, so the run was repeated.

Real occlusion against the dense scan. Against all 49 DTU views (Figure 10), 1.1 to 1.7% of the predicted points of the four backbones lie on surfaces no input view saw with three views, 0.2 to 0.4% with ten. On ETH3D, where each scene has few images, the share is 3.9 to 8.0%. Where the share allows a claim, the head beats native confidence in 10 of 14 settings (Gates 8 and 10): by 0.09 to 0.13 AUROC for every backbone on DTU at five views and for Pi3X, Pi3 and DA3 at three, and by 0.07 to 0.08 for VGGT on ETH3D at every view count. It does not for VGGT on DTU at three views (+0.02) or for Pi3X on ETH3D (+0.00 to +0.03). The models rarely invent geometry their inputs give no reason for; where they do, the head finds it more often than not.

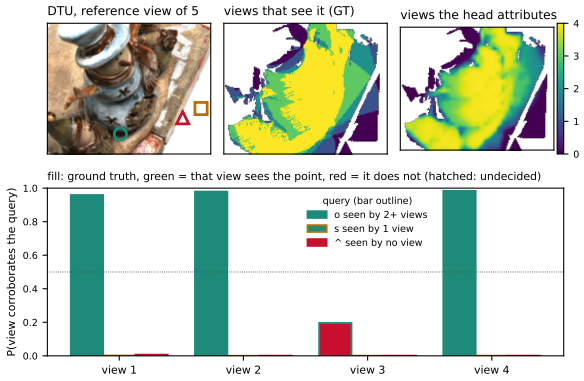


Figure 12. **Which view saw it** (provenance variant, VGGT, an unseen DTU scan). Top: views that see each pixel and views the head attributes. Bottom: the head's probability per view for three query pixels; bar fill is the ground truth.

Which view saw it. The provenance variant (Figure 12) ranks the views that see a point above those that do not at prov-AUROC 0.93 to 0.96 on DTU and 0.83 to 0.88 on CO3Dv2 with no camera, the level of reprojection there, and trails reprojection on building scans and landmarks, so the convolutional head stays the method of record (Gate 5, supplement).

Adapting without labels. Trained label-free on half of a target dataset and scored on the other half, the head gains where the backbone's own geometry is reliable (Pi3X on CO3Dv2 and ETH3D, +0.025 and +0.027) and loses where its teacher is weak (VGGT on CO3Dv2, −0.09); the DTU-trained head is the safe default.

Downstream: a certified reconstruction. The head keeps almost no single-view geometry until high coverage (Figure 11a). With the bands fixed in advance, the head marks 57% to 87% of the pixels *observed*, of which 0.5% to 5% are single-view; a threshold set on other scenes lets 1% to 2% of single-view pixels through on DTU against 57% to 83% for native confidence. Keeping the 80% of points the head trusts most raises the accuracy of the fused DTU cloud from 0.65 to 0.77 (Figure 11b); at that coverage the head keeps half as much single-view geometry as native confidence (2.5% against 4.6% of the points) at nearly the same accuracy (0.77 against 0.78). Abstaining thus yields a cloud that is as accurate and twice as well corroborated. The two scores measure different things: native confidence ranks depth error, the head ranks support, and averaging them changes F1 by less than 0.01 (Gate 11). A user can apply one filter for accuracy and the other to certify what was seen.

The limit. No cross-view test fires on a surface every view shows (Figure 13). On the benchmark's own patches the visibility heads score 0.70 to 0.75 AUROC, and disagreement between two different backbones 0.52 to 0.61: five models

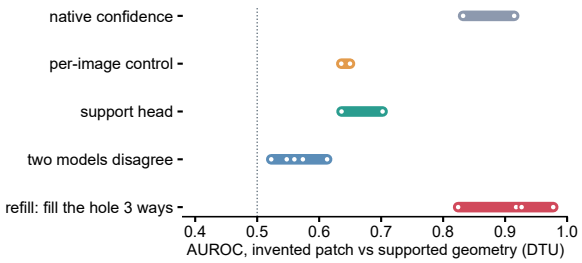


Figure 13. **The limit.** On the benchmark's own patch, which every view shows, no cross-view test separates the invented surface; refilling the hole does. Native confidence scores high only because it reacts to the patch (Table 1).

Table 3. **Pre-registered pass conditions,** each written with a date and a commit before the run it judges (supplement). ✓ pass, ● partial, ✗ fail.

gate	condition	outcome
2	fill-trained head beats its per-image control	✗
3	ground-truth head beats native confidence and the per-image control (margin 0.02), on DTU and an unseen dataset	✓
4	label-free head: as Gate 3, within 0.03 of the ground-truth head, ahead of its reprojection teacher on CO3Dv2	✓
5	provenance head beats reprojection attribution on three datasets	✗
6	native confidence still leaks under the matched control	✓
7	matched head, four backbones: control at most 0.60, head beats it by 0.10 and native confidence; 0.75 with one low-parallax view	●
8	real occlusion on DTU: share at least 0.5% and head beats native confidence	●
9	a one-view fine-tune reaches 0.75 with one low-parallax view	✗
10	real occlusion with four backbones and on ETH3D: as Gate 8 (holds in 10 of 14 settings)	●
11	native confidence and head combined improve fused-cloud F1	✗

of three families invent the same surface. That surface is not fixed by the prior, though. Filling the same hole three ways moves it, and the spread separates the patch at AUROC 0.82 to 0.98 on DTU: the model reads the hole and draws what its pixels suggest.

Pre-registration. Table 3 lists every pass condition and its outcome, failures included. Gate 2 passed by the letter on a 0.003 difference and is not counted; it is why every later gate carries a margin and a per-image control.

Takeaway. The backbone knows which of its geometry has multi-view support. A small head reads it without labels; what no cross-view test can read is an invention every view agrees on.

7. Limitations and conclusion

Limitations. The head certifies multi-view agreement, not truth: a surface every view agrees on passes, whether a camera saw it or the prior drew it. The benchmark's labels are a proxy for real occlusion, which the stress tests, the matched split and the dense-scan comparisons on DTU and ETH3D bound but do not remove. CO3Dv2 depth is COLMAP with 1% to 3% noise and three backbones saw CO3Dv2 in training; the landmarks' ground truth has holes; scenes are static. A single low-parallax view remains the hardest matched case, and the matched split was run on DTU and CO3Dv2.

Conclusion. Feed-forward 3D models mostly do not say when they have not seen a surface, and their confidence reacts to appearance. Their cross-view tokens do carry the answer, and a small head trained on the backbone's own agreement reads it on captures it never saw. The next problem is the invention every view agrees on.

LLM Usage Statement

A general-purpose LLM assistant was used, under the authors' direction, for parts of the code (the benchmark runner, the model wrappers, the run scripts and figure scripts), for literature lookup, and for editing the text of this paper, including grammar. The study design, the definitions and metrics, the choice and freezing of every hyperparameter, the execution and checking of every experiment, the verification of every reported number against its logged run, and the conclusions are the authors' own, and the authors are responsible for all content. The assistant was not used to generate results, figures from unlogged data, or references; every citation was checked against the published record. The tool is named in the camera-ready version.

References

[1] David Ahmedt-Aristizabal, Mohammad Ali Armin, Russell Tsuchida, and Lars Petersson. GeoCond: A Conditioning-Aware reliability adapter for Feed-Forward 3D reconstruction. *arXiv:2609.18465*, 2026. 2

[2] Denis M. Akola and David F. Fouhey. Zero-Shot novel depth synthesis using 3D foundation models scene representations. *arXiv:2609.04174*, 2026. 2

[3] Jelena Bratulić, Sudhanshu Mittal, Thomas Brox, and Christian Rupprecht. On geometric understanding and learned priors in Feed-forward 3D reconstruction models. *arXiv:2512.11508*, 2025. 2

[4] Jinhyeok Choi, JoungBin Lee, Jongmin Lee, Minkyeong Jeon, Jiyoung Kim, Kaeun Han, Yoonjin Yoon, and Seungryong Kim. L3G-World: Real-time infinite world generation with latent 3D gaussian memory. *arXiv preprint arXiv:2610.40614*, 2026. 2

[5] Alibek Dadajonov and Hikmatillo Umarhojiyev. Sharper frames, worse rooms: Viewpoint coverage dominates image

quality in feed-forward 3D reconstruction. *arXiv preprint arXiv:2610.15854*, 2026. 2

[6] Haoge Deng, Ting Pan, Fan Zhang, Yang Liu, Zhuoyan Luo, Yufeng Cui, Wenxuan Wang, Chunhua Shen, Shiguang Shan, Zhaoxiang Zhang, and Xinlong Wang. Uniform discrete diffusion with metric path for video generation. *arXiv:2510.24717*, 2025. URSA. 2

[7] Nanxing Nick Deng, Qing Cheng, Niclas Zeller, and Daniel Cremers. A calibration audit of confidence in Feed-Forward 3D reconstruction models. *arXiv:2608.29705*, 2026. 2

[8] Rasmus Jensen, Anders Dahl, George Vogiatzis, Engin Tola, and Henrik Aanæs. Large scale multi-view stereopsis evaluation. In *CVPR*, 2014. 2

[9] Lihan Jiang, Yucheng Mao, Linning Xu, Tao Lu, Kerui Ren, Yichen Jin, Xudong Xu, Mulin Yu, Jiangmiao Pang, Feng Zhao, Dahua Lin, and Bo Dai. AnySplat: Feed-forward 3D gaussian splatting from unconstrained views. *arXiv preprint arXiv:2505.23716*, 2025. 2, 3

[10] Yuhe Jin, Dmytro Mishkin, Anastasiia Mishchuk, Jiří Matas, Pascal Fua, Kwang Moo Yi, and Eduard Trulls. Image matching across wide baselines: From paper to practice. *International Journal of Computer Vision*, 129:517–547, 2021. 3

[11] Jinyoung Jun, Lei Chu, Jiahao Li, Yan Lu, and Chang-Su Kim. Perfecting depth: Uncertainty-Aware enhancement of metric depth. *arXiv:2506.04612*, 2025. 2

[12] Bingxin Ke, Anton Obukhov, Shengyu Huang, Nando Metzger, Rodrigo Caye Daudt, and Konrad Schindler. Repurposing Diffusion-Based image generators for Monocular Depth estimation. In *CVPR*, 2024. 2

[13] Nikhil Keetha, Norman Müller, Johannes Schönberger, Lorenzo Porzi, Yuchen Zhang, Tobias Fischer, Arno Knapitsch, Duncan Zauss, Ethan Weber, Nelson Antunes, Jonathon Luiten, Manuel Lopez-Antequera, Samuel Rota Bulò, Christian Richardt, Deva Ramanan, Sebastian Scherer, and Peter Kontschieder. MapAnything: Universal Feed-Forward metric 3D reconstruction. *arXiv:2509.13414*, 2025. 1, 2, 3

[14] Vincent Leroy, Yohann Cabon, and Jerome Revaud. Grounding image matching in 3D with MASt3R. In *ECCV*, 2024. 1, 2

[15] Daisy Li, Kyle Gao, Quanyun Wu, Hanna Chomko, John S. Zelek, and Jonathan Li. When wider views fail: Stress-testing feed-forward 3D reconstruction. *arXiv preprint arXiv:2609.24839*, 2026. 2

[16] Yifan Li, Yifan Du, Kun Zhou, Jinpeng Wang, Wayne Xin Zhao, and Ji-Rong Wen. Evaluating object hallucination in large Vision-Language models. In *EMNLP*, 2023. 2

[17] Haotong Lin, Sili Chen, Junhao Liew, Donny Y. Chen, Zhenyu Li, Guang Shi, Jiashi Feng, and Bingyi Kang. Depth Anything 3: Recovering the visual space from any views. *arXiv:2511.10647*, 2025. 1, 2, 3

[18] Sammam Mahdi, Fariha Binta Salim, Rakin Bin Rabbani, and Aniqua Nusrat Zereen. ReVar3R: Gauge-Aware perturbation uncertainty for Feed-Forward 3D reconstruction. *arXiv:2610.07883*, 2026. 2

[19] Fabian Mentzer, David Minnen, Eirikur Agustsson, and Michael Tschannen. Finite scalar quantization: VQ-VAE made simple. In *ICLR*, 2024. 2

[20] Tuan Duc Ngo, Chuang Gan, and Evangelos Kalogerakis. VolFill: Single-View amodal 3D scene reconstruction with volumetric flow matching. *arXiv:2605.31466*, 2026. 2

[21] Hoang Nguyen, Xiaohao Xu, and Xiaonan Huang. The 3D mirage: Probing and taming 3D hallucinations. *arXiv:2512.15423*, 2025. 2

[22] Jingyi Pan, Dan Xu, and Qiong Luo. Pose-free feed-forward 3D inpainting via learnable mask attention and support token refinement. *arXiv preprint arXiv:2610.11857*, 2026. 2

[23] Soumava Paul, Prakhar Kaushik, and Alan Yuille. Can these views be one scene? Evaluating multiview 3D consistency when 3D foundation models hallucinate. *arXiv:2605.18754*, 2026. 2

[24] Jeremy Reizenstein, Roman Shapovalov, Philipp Henzler, Luca Sbordone, Patrick Labatut, and David Novotny. Common objects in 3D: Large-Scale learning and evaluation of real-life 3D category reconstruction. In *ICCV*, 2021. 2

[25] Thomas Schöps, Johannes L. Schönberger, Silvano Galliani, Torsten Sattler, Konrad Schindler, Marc Pollefeys, and Andreas Geiger. A multi-view stereo benchmark with high-resolution images and multi-camera videos. In *CVPR*, 2017. 2

[26] Giuseppe Stracquadanio, Kevin Raj, Julia Grabinski, and Stefan Roth. ReconSplat: Generalizable 3D scene reconstruction beyond observed views. *arXiv preprint arXiv:2608.28895*, 2026. 2

[27] Alexandru Telea. An image inpainting technique based on the Fast Marching Method. *Journal of Graphics Tools*, 9(1), 2004. 2

[28] Aditya Thirumani Srinivasan and David Nordström. Does VGGT know what it does not know? *arXiv preprint arXiv:2610.35806*, 2026. 2

[29] Jianyuan Wang, Minghao Chen, Nikita Karaev, Andrea Vedaldi, Christian Rupprecht, and David Novotny. VGGT: Visual geometry grounded transformer. In *CVPR*, 2025. 1, 2, 3

[30] Jianyuan Wang, Minghao Chen, Shangzhan Zhang, Nikita Karaev, Johannes Schönberger, Patrick Labatut, Piotr Bojanowski, David Novotny, Andrea Vedaldi, and Christian Rupprecht. VGGT-ω. *arXiv:2605.15195*, 2026. 2

[31] Shuzhe Wang, Vincent Leroy, Yohann Cabon, Boris Chidlovskii, and Jerome Revaud. DUSt3R: Geometric 3D vision made easy. In *CVPR*, 2024. 1, 2

[32] Yifan Wang, Jianjun Zhou, Haoyi Zhu, Wenzheng Chang, Yang Zhou, Zizun Li, Junyi Chen, Jiangmiao Pang, Chunhua Shen, and Tong He. π^3: Scalable Permutation-Equivariant visual geometry learning. *arXiv:2507.13347*, 2025. 1, 2, 3

[33] Zihao Zhu, Wenyuan Zhao, Nuo Chen, Chao Tian, and Zhiwen Fan. Trust it or not: Evidential uncertainty for Feed-Forward 3D reconstruction with Trust3R. *arXiv:2605.19539*, 2026. 2

[34] Filippo Ziliotto, Luciano Serafini, Lamberto Ballan, and Tommaso Campari. What VGGT knows about overlap: Probing geometric foundation models for co-visibility. *arXiv preprint arXiv:2607.09503*, 2026. 2